

**AA 203**

# **Optimal and Learning-Based Control**

Course overview; Feedback, stability, and optimal control problems

Autonomous Systems Laboratory

Stanford University

April 3, 2023

(last updated April 12, 2023)



**Stanford**  
University

## Teaching team



Spencer M. Richards  
(Instructor)



Daniele Gammelli  
(Instructor)



Thomas Lew  
(CA)



Devansh Jalota  
(CA)

## Important links

- Lecture slides and homework assignments:  
<https://asl.stanford.edu/aa203>
- Lecture recordings:  
<https://canvas.stanford.edu/courses/171491>
- Announcements and discussion forum:  
<https://edstem.org/us/courses/38294>
- Coursework submission:  
<https://www.gradescope.com/courses/525712>
- For urgent questions:  
[aa203-spr2223-staff@lists.stanford.edu](mailto:aa203-spr2223-staff@lists.stanford.edu)

## Homework (60%)

- 4 homeworks, each worth 15%.
- Covers a mixture of theory and programming.
- Generally due every 2 weeks.

## Project (40%)

- 5% proposal, 10% midterm report, 25% final report and video presentation.
- Open-ended in groups of up to 3 people.

## Discussion ( $\leq 5\%$ bonus)

- 0.5% per endorsed Edstem post, up to 5%.

## Late days

- 6 total, up to 3 on a single assignment.
- *Not applicable* to the final report and video presentation (due on the last day of class).

In order of importance:

**Lecture slides** Should be posted on the class website before each lecture.

**Recitations** Friday lecture sessions (Weeks 1–4) led by the CAs covering supplementary tools (mathematical and computational).

**Course notes** Evolving, somewhat outdated partial notes available at:  
<https://github.com/StanfordASL/AA203-Notes>

**Textbooks** Suggested ad hoc during lecture and discussions (not required).

# Prerequisites

- Standard undergraduate engineering mathematics knowledge (i.e., vector calculus, ordinary differential equations (ODEs), probability theory).
- *Strong* familiarity with linear algebra (e.g., EE263, CME200).
- Some knowledge of optimization is nice to have (e.g., EE364A, CME307, CS269O, AA222).
- To get the most out of this class, it is recommended to have taken at least one course in:
  - control (e.g., ENGR105, ENGR205, AA212)
  - machine learning (e.g., CS229, CS230, CS231N)
- *Homework 0 (ungraded)* is out now to help you gauge your preparedness.

- Arguably, this class aims for breadth over depth. Some past students have needed to self-study some of the details.
- The course content is subject to feedback. Homework problems covering state-of-the-art topics sometimes suffer from bugs.
- This class is quite challenging. Some past students have had trouble managing both homeworks and project deliverables.
- Projects focused on *learning-based* control may require self-study of material before the relevant lectures.

# Agenda


1. Context and course goals
2. Stability and Lyapunov functions
3. Optimal control problems



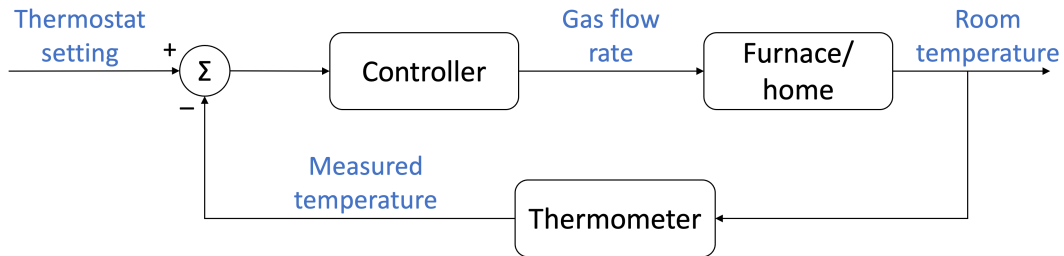
# Agenda

1. Context and course goals
2. Stability and Lyapunov functions
3. Optimal control problems

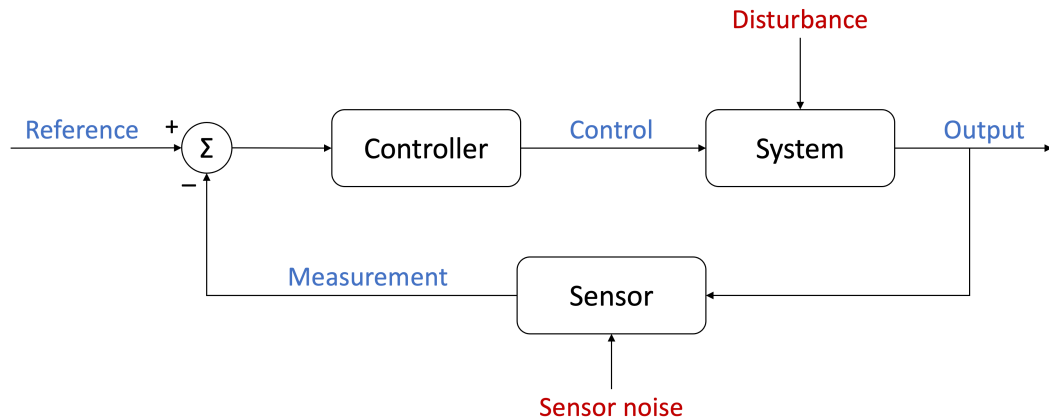
# Feedback control




## Feedback control example



## Feedback control with complications



## Feedback control in reinforcement learning



Continuous-time:

Time  $t \in \mathbb{R}$

State  $x(t) \in \mathbb{R}^n$

Control input  $u(t) \in \mathbb{R}^m$

Dynamics  $\dot{x}(t) = f(t, x(t), u(t))$

Trajectories  $x : t \mapsto x(t)$   
 $u : t \mapsto u(t)$

Discrete-time:

$t \in \mathbb{N}$

$x_t \in \mathbb{R}^n$

$u_t \in \mathbb{R}^m$

$x_{t+1} = f(t, x_t, u_t)$

$x : t \mapsto x_t$   
 $u : t \mapsto u_t$

We assume  $f$  is sufficiently “well-behaved” such that, given a piecewise-continuous input  $u$ , there exists a unique solution  $x$  for each initial condition.

In roughly the second-half of the course, the dynamics may be *unknown*, and so will have to *learn* how to control our system based on data.

## Example: Double-integrator control

Point-mass with acceleration control in 1-D:

$$\begin{pmatrix} \dot{s} \\ \dot{v} \end{pmatrix} = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{pmatrix} s \\ v \end{pmatrix} + \begin{bmatrix} 0 \\ 1 \end{bmatrix} u$$



More generally, in multiple dimensions we have:

$$\begin{pmatrix} \dot{s} \\ \dot{v} \end{pmatrix} = \begin{bmatrix} 0 & I \\ 0 & 0 \end{bmatrix} \begin{pmatrix} s \\ v \end{pmatrix} + \begin{bmatrix} 0 \\ I \end{bmatrix} u$$

**Objective** Drive to a standstill at the origin, i.e.,  $(0,0)$ .

**Proposal** Proportional-derivative (PD) feedback:

$$u = -k_p s - k_d v \implies \begin{pmatrix} \dot{s} \\ \dot{v} \end{pmatrix} = \begin{bmatrix} 0 & 1 \\ -k_p & -k_d \end{bmatrix} \begin{pmatrix} s \\ v \end{pmatrix}$$



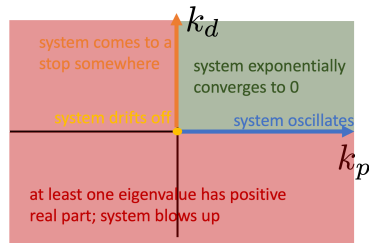
## Example: Double-integrator stability

Is the closed-loop system stable?

$$\begin{pmatrix} \dot{s} \\ \dot{v} \end{pmatrix} = \begin{bmatrix} 0 & 1 \\ -k_p & -k_d \end{bmatrix} \begin{pmatrix} s \\ v \end{pmatrix} \Rightarrow \begin{pmatrix} s(t) \\ v(t) \end{pmatrix} = \underbrace{\exp\left(\begin{bmatrix} 0 & 1 \\ -k_p & -k_d \end{bmatrix} t\right)}_{=:\Phi(t)} \begin{pmatrix} s(0) \\ v(0) \end{pmatrix}$$

where  $\Phi(t) = V \exp(tJ) V^{-1}$  with eigenvalues  $\lambda_{\pm} = -\frac{k_d}{2} \pm \frac{1}{2} \sqrt{k_d^2 - 4k_p}$  and

$$\exp(tJ) = \begin{cases} \begin{bmatrix} e^{\lambda_+ t} & 0 \\ 0 & e^{\lambda_- t} \end{bmatrix}, & k_d^2 > 4k_p \\ \begin{bmatrix} 1 & t \\ 0 & 1 \end{bmatrix} e^{-\frac{k_d}{2} t}, & k_d^2 = 4k_p \\ \begin{bmatrix} \cos(\omega t) & \sin(\omega t) \\ -\sin(\omega t) & \cos(\omega t) \end{bmatrix} e^{-\frac{k_d}{2} t}, & k_d^2 < 4k_p \end{cases}$$





Traditional feedback control balances the following desiderata.

**Stability** The system output does not diverge or “blow up”.

**Tracking** The system output converges to a desired reference.

**Disturbance rejection** The system is insensitive to disturbances and noise.

**Robustness** The controller performs well despite some model misspecification.

This course also incorporates and focuses on the following objectives.

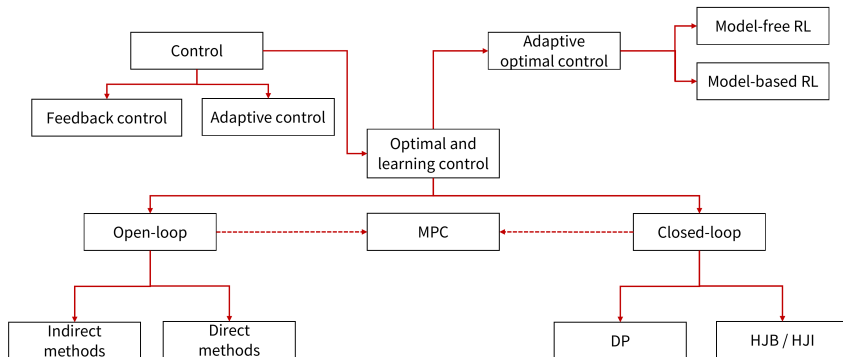
**Performance** The controller achieves an optimal trade-off between various metrics.

**Constraints** The controller does not cause the system to violate safety restrictions or inherent (e.g., physical) limitations.

**Planning** An appropriate reference trajectory is computed and given to the controller for tracking.

**Learning** The controller can adapt to an unknown or time-varying system.

# Course overview and goals



- To learn the *theory* and *practice* of fundamental techniques in optimal and learning-based control.
- To gain a *holistic understanding* of how such techniques are used across fields.

# Agenda

1. Context and course goals
2. Stability and Lyapunov functions
3. Optimal control problems

## Mathematical definitions of stability

Consider  $\dot{x} = f(x)$  (or  $\dot{x} = f(x, \pi(x))$ ) and an *equilibrium*  $\bar{x} \in \mathbb{R}^n$  (i.e.,  $f(\bar{x}) = 0$ ).

**Marginal/Lyapunov**  $\forall \varepsilon > 0, \exists \delta > 0 : \|x(0) - \bar{x}\| < \delta \implies \|x(t) - \bar{x}\| < \varepsilon, \forall t \geq 0$   
“Trajectories that start close to the equilibrium remain close to the equilibrium.”

**Asymptotic (local)**  $\exists \delta > 0 : \|x(0) - \bar{x}\| < \delta \implies \lim_{t \rightarrow \infty} \|x(t) - \bar{x}\| = 0$   
“Trajectories that start near the equilibrium converge to it.”

**Exponential (local)**  $\exists \delta, c, \alpha > 0 : \|x(0) - \bar{x}\| < \delta \implies \|x(t) - \bar{x}\| \leq ce^{-\alpha t} \|x(0) - \bar{x}\|$   
“Trajectories that start near the equilibrium converge to it exponentially fast.”

Take  $\delta \rightarrow \infty$  to get “global” definitions. For linear time-invariant (LTI) systems, “asymptotic = exponential” and “local = global” always.


## Theorem (Lyapunov's direct method)

Consider  $\dot{x} = f(x)$  where  $f$  is locally Lipschitz and  $f(0) = 0$ . Suppose there exists  $V \in \mathcal{C}^1(\mathbb{R}^n, \mathbb{R})$  such that

- $V$  is positive-definite, i.e.,  $V(x) \geq 0$  and  $V(x) = 0 \iff x = 0$ ,
- $\dot{V}$  is negative-definite, i.e.,  $\nabla V(x)^\top f(x) \leq 0$  and  $\nabla V(x)^\top f(x) = 0 \iff x = 0$ .

Then  $\bar{x} = 0$  is locally asymptotically stable. If in addition

- $V$  is radially unbounded, i.e.,  $V(x) \rightarrow \infty$  as  $\|x\| \rightarrow \infty$ , then  $\bar{x} = 0$  is globally asymptotically stable.



If the “energy”  $V(x)$  is decreasing everywhere along trajectories, then  $V(x) \rightarrow 0$  and thus  $x \rightarrow 0$ .

The existence of a Lyapunov function is a sufficient condition or *certificate* for stability. Pointwise Lyapunov inequalities are generally less cumbersome to work with than limits.

## Converse Lyapunov theorems

The existence of a Lyapunov function is also necessary for stability.

### Theorem (Converse Lyapunov theorem)


Consider  $\dot{x} = f(x)$  where  $f$  is locally Lipschitz. Suppose  $\bar{x} = 0$  is a locally asymptotically stable equilibrium with region of attraction  $\mathcal{A} \subset \mathbb{R}^n$ . Then there exists

$V \in \mathcal{C}^1(\mathbb{R}^n, \mathbb{R})$  such that

- $V$  is positive-definite on  $\mathcal{A}$ ,
- $\dot{V}$  is negative-definite on  $\mathcal{A}$ ,
- $V(x) \rightarrow \infty$  as  $x \rightarrow \partial\mathcal{A}$  (the boundary of  $\mathcal{A}$ ),
- $\{x \mid V(x) \leq c\}$  is a compact subset of  $\mathcal{A}$  for any  $c > 0$ .

If  $\bar{x} = 0$  is globally asymptotically stable, i.e.,  $\mathcal{A} = \mathbb{R}^n$ , then

- $V$  is radially unbounded.



If the “energy”  $V(x)$  is decreasing everywhere along trajectories, then  $V(x) \rightarrow 0$  and thus  $x \rightarrow 0$ .

# Agenda

1. Context and course goals
2. Stability and Lyapunov functions
3. Optimal control problems



## Optimal control problems (continuous-time)

$$\underset{x,u}{\text{minimize}} \quad J(x,u) := \ell_T(T, x(T)) + \int_0^T \ell(t, x(t), u(t)) dt \quad \text{cost (terminal + stage)}$$

$$\text{subject to} \quad \dot{x}(t) = f(t, x(t), u(t)), \quad \forall t \in [0, T] \quad \text{dynamical feasibility}$$

$$x(t_0) = x_0, \quad x(T) \in \mathcal{X}_T \quad \text{boundary conditions}$$

$$x(t) \in \mathcal{X}, \quad \forall t \in [0, T] \quad \text{state constraints}$$

$$u(t) \in \mathcal{U}, \quad \forall t \in [0, T] \quad \text{input constraints}$$

An optimal control  $u^*(t)$  for a specific initial state  $x_0$  is an *open-loop* input. An optimal control of the form  $u^*(t) = \pi^*(t, x(t))$  is a *closed-loop* input.

The stochastic and unknown model settings will be covered later on in the course.

## Optimal control problems (discrete-time)

$$\begin{aligned} \underset{x,u}{\text{minimize}} \quad & J(x,u) := \ell_T(T, x_T) + \sum_{t=0}^{T-1} \ell(t, x_t, u_t) && \text{cost (terminal + stage)} \\ \text{subject to} \quad & x_{t+1} = f(t, x_t, u_t), \quad \forall t \in \{0, 1, \dots, T-1\} && \text{dynamical feasibility} \\ & x_0 = \bar{x}_0, \quad x_T \in \mathcal{X}_T && \text{boundary conditions} \\ & x_t \in \mathcal{X}, \quad \forall t \in \{0, 1, \dots, T-1\} && \text{state constraints} \\ & u_t \in \mathcal{U}, \quad \forall t \in \{0, 1, \dots, T-1\} && \text{input constraints} \end{aligned}$$

An optimal control  $u_t^*$  for a specific initial state  $x_0$  is an *open-loop* input. An optimal control of the form  $u_t^* = \pi^*(t, x_t)$  is a *closed-loop* input.

The stochastic and unknown model settings will be covered later on in the course.

## Example: Finite-horizon linear quadratic regulator (LQR)

$$\underset{x,u}{\text{minimize}} \quad x(T)^\top Q_T x(T) + \int_0^T \left( x(t)^\top Q(t) x(t) + u(t)^\top R(t) u(t) \right) dt \quad \text{cost}$$

$$\text{subject to} \quad \dot{x}(t) = A(t)x(t) + B(t)u(t), \quad \forall t \in [0, T]$$

dynamical feasibility

$$x(0) = x_0$$

initial condition

For linear dynamics and a quadratic cost, we can derive *the* optimal feedback law  $u^*(t) = K(t)x(t)$ , which is also linear.

## Example: Infinite-horizon linear quadratic regulator (LQR)

$$\underset{x,u}{\text{minimize}} \quad \int_0^\infty \left( x(t)^\top Q x(t) + u(t)^\top R u(t) \right) dt \quad \text{cost}$$

$$\begin{aligned} \text{subject to } \dot{x}(t) &= Ax(t) + Bu(t), \quad \forall t \in [0, \infty) && \text{dynamical feasibility} \\ x(t_0) &= x_0 && \text{initial condition} \end{aligned}$$

For LTI dynamics and a time-invariant quadratic cost, we can derive *the* optimal feedback law  $u^*(t) = Kx(t)$ , which is also LTI.

The closed-loop system must converge to zero (i.e., be asymptotically stable) to ensure the infinite-horizon cost is well-defined.

The cost function  $J(x^*, u^*)$  is a Lyapunov function for the closed-loop dynamics!

Nonlinear optimization theory  
(for unconstrained and constrained problems)