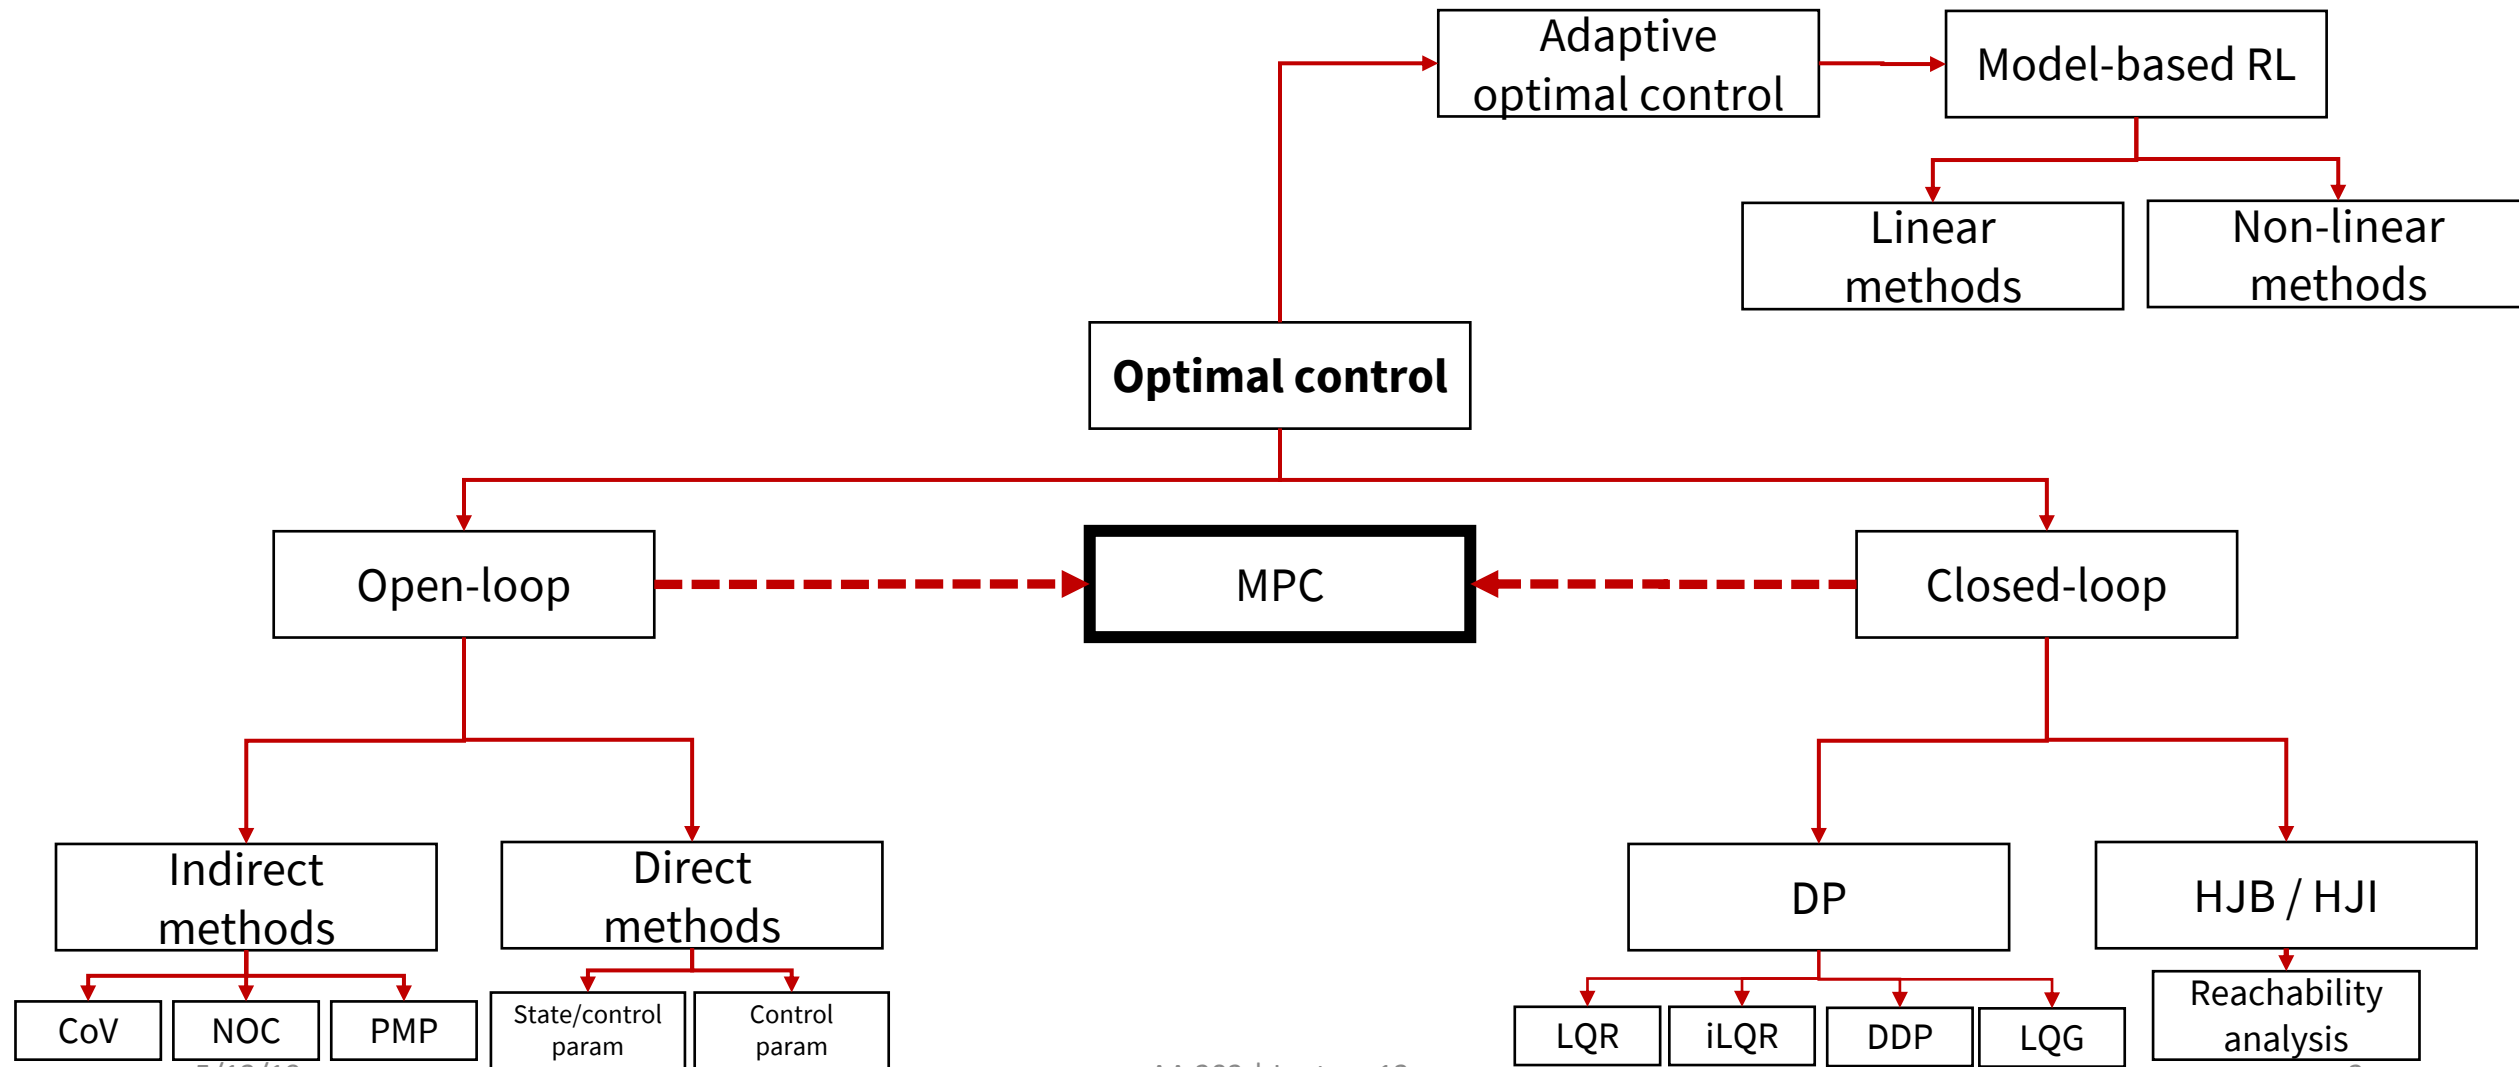


AA203

Optimal and Learning-based Control


Introduction to MPC, persistent feasibility

Roadmap



Model predictive control

- Model predictive control (or, more broadly, receding horizon control) entails solving finite-time optimal control problems in a receding horizon fashion



Model predictive control

Key steps:

1. At each sampling time t , solve an *open-loop* optimal control problem over a finite horizon
2. Apply optimal input signal during the following sampling interval $[t, t + 1)$
3. At the next time step $t + 1$, solve new optimal control problem based on new measurements of the state over a shifted horizon

Basic formulation

- Consider the problem of regulating to the origin the discrete-time linear invariant system

$$\mathbf{x}(t + 1) = A\mathbf{x}(t) + B\mathbf{u}(t), \quad \mathbf{x}(t) \in \mathbb{R}^n, \quad \mathbf{u}(t) \in \mathbb{R}^m$$

subject to the constraints

$$\mathbf{x}(t) \in X, \quad \mathbf{u}(t) \in U, \quad t \geq 0$$

where the sets X and U are *polyhedra*

Basic formulation

- Assume that a full measurement of the state $\mathbf{x}(t)$ is available at the current time t
- The finite-time optimal control problem solved at each stage is

$$J_t^*(\mathbf{x}(t)) = \min_{\mathbf{u}_{t|t}, \dots, \mathbf{u}_{t+N-1|t}} p(\mathbf{x}_{t+N|t}) + \sum_{k=0}^{N-1} c(\mathbf{x}_{t+k|t}, \mathbf{u}_{t+k|t})$$

$$\text{subject to } \mathbf{x}_{t+k+1|t} = A\mathbf{x}_{t+k|t} + B\mathbf{u}_{t+k|t}, \quad k = 0, \dots, N-1$$

$$\mathbf{x}_{t+k|t} \in X, \quad \mathbf{u}_{t+k|t} \in U, \quad k = 0, \dots, N-1$$

$$\mathbf{x}_{t+N|t} \in X_f$$

$$\mathbf{x}_{t|t} = \mathbf{x}(t)$$

Basic formulation

Notation:

- $\mathbf{x}_{t+k|t}$ is the state vector at time $t + k$ predicted at time t (via the system's dynamics)
- $\mathbf{u}_{t+k|t}$ is the input \mathbf{u} at time $t + k$ computed at time t

Note: $\mathbf{x}_{3|1} \neq \mathbf{x}_{3|2}$

Basic formulation

- Let $U_{t \rightarrow t+N|t}^* := \{\mathbf{u}_{t|t}^*, \mathbf{u}_{t+1|t}^*, \dots, \mathbf{u}_{t+N-1|t}^*\}$ be the optimal solution, then

$$\mathbf{u}(t) = \mathbf{u}_{t|t}^*(\mathbf{x}(t))$$

- The optimization problem is then repeated at time $t + 1$, based on the new state $\mathbf{x}_{t+1|t+1} = \mathbf{x}(t + 1)$
- Define $\pi_t(\mathbf{x}(t)) := \mathbf{u}_{t|t}^*(\mathbf{x}(t))$
- Then the closed-loop system evolves as
$$\mathbf{x}(t + 1) = A\mathbf{x}(t) + B\pi_t(\mathbf{x}(t)) := f_{cl}(\mathbf{x}(t), t)$$
- Central question: characterize the behavior of **closed-loop** system

Simplifying the notation

- Note that the setup is time-invariant, hence, to simplify the notation, we can let $t = 0$ in the finite-time optimal control problem, namely

$$\begin{aligned} J_0^*(\mathbf{x}(t)) &= \min_{\mathbf{u}_0, \dots, \mathbf{u}_{N-1}} p(\mathbf{x}_N) + \sum_{k=0}^{N-1} c(\mathbf{x}_k, \mathbf{u}_k) \\ \text{subject to } \quad &\mathbf{x}_{k+1} = A\mathbf{x}_k + B\mathbf{u}_k, \quad k = 0, \dots, N-1 \\ &\mathbf{x}_k \in X, \quad \mathbf{u}_k \in U, \quad k = 0, \dots, N-1 \\ &\mathbf{x}_N \in X_f \\ &\mathbf{x}_0 = \mathbf{x}(t) \end{aligned}$$

- Denote $U_0^*(\mathbf{x}(t)) = \{\mathbf{u}_0^*, \dots, \mathbf{u}_{N-1}^*\}$

Simplifying the notation

- With new notation,

$$\mathbf{u}(t) = \mathbf{u}_0^*(\mathbf{x}(t)) = \pi(\mathbf{x}(t))$$

and closed-loop system becomes

$$\mathbf{x}(t+1) = A\mathbf{x}(t) + B\pi(\mathbf{x}(t)) := f_{cl}(\mathbf{x}(t))$$

Typical cost functions

- 2-norm:

$$p(\mathbf{x}_N) = \mathbf{x}_N' P \mathbf{x}_N, \quad c(\mathbf{x}_k, \mathbf{u}_k) = \mathbf{x}_k' Q \mathbf{x}_k + \mathbf{u}_k' R \mathbf{u}_k, \quad P \geq 0, Q \geq 0, R > 0$$

- 1-norm or ∞ -norm:

$$p(\mathbf{x}_N) = \|P \mathbf{x}_N\|_p \quad c(\mathbf{x}_k, \mathbf{u}_k) = \|Q \mathbf{x}_k\|_p + \|R \mathbf{u}_k\|_p, \quad p = 1 \text{ or } \infty$$

where P, Q, R are full column ranks

Online model predictive control

repeat

measure the state $\mathbf{x}(t)$ at time instant t

obtain $U_0^*(\mathbf{x}(t))$ by solving finite-time optimal control problem

if $U_0^*(\mathbf{x}(t)) = \emptyset$ **then** ‘problem infeasible’ **stop**

apply the first element \mathbf{u}_0^* of $U_0^*(\mathbf{x}(t))$ to the system

wait for the new sampling time $t + 1$

Main implementation issues

1. The controller may lead us into a situation where after a few steps the finite-time optimal control problem is infeasible → *persistent feasibility issue*
2. Even if the feasibility problem does not occur, the generated control inputs may not lead to trajectories that converge to the origin (i.e., closed-loop system is unstable) → *stability issue*

Key question: how do we guarantee that such a “short- sighted” strategy leads to effective long-term behavior?

Analysis approaches

1. Analyze closed-loop behavior directly → generally very difficult
2. Derive conditions on terminal function p , and terminal constraint set X_f so that persistent feasibility and closed-loop stability are guaranteed

Addressing persistent feasibility

Goal: design MPC controller so that feasibility for all future times is guaranteed

Approach: leverage tools from *invariant set theory*

Set of feasible initial states

- Set of feasible initial states

$$X_0 := \{\mathbf{x}_0 \in X \mid \exists (\mathbf{u}_0, \dots, \mathbf{u}_{N-1}) \text{ such that } \mathbf{x}_k \in X, \mathbf{u}_k \in U, k = 0, \dots, N-1, \\ \mathbf{x}_N \in X_f \text{ where } \mathbf{x}_{k+1} = A\mathbf{x}_k + B\mathbf{u}_k, k = 0, \dots, N-1\}$$

- A control input can be found only if $\mathbf{x}(0) \in X_0$!

Controllable sets

- For the autonomous system $\mathbf{x}(t + 1) = \phi(\mathbf{x}(t))$ with constraints $\mathbf{x}(t) \in X, \mathbf{u}(t) \in U$, the one-step controllable set to set S is defined as

$$\text{Pre}(S) := \{\mathbf{x} \in \mathbb{R}^n : \phi(\mathbf{x}) \in S\}$$

- For the system $\mathbf{x}(t + 1) = \phi(\mathbf{x}(t), \mathbf{u}(t))$ with constraints $\mathbf{x}(t) \in X, \mathbf{u}(t) \in U$, the one-step controllable set to set S is defined as
- $$\text{Pre}(S) := \{\mathbf{x} \in \mathbb{R}^n : \exists \mathbf{u} \in U \text{ such that } \phi(\mathbf{x}, \mathbf{u}) \in S\}$$

Control invariant sets

- A set $\mathcal{C} \subseteq X$ is said to be a control invariant set for the system $\mathbf{x}(t+1) = \phi(\mathbf{x}(t), \mathbf{u}(t))$ with constraints $\mathbf{x}(t) \in X$, $\mathbf{u}(t) \in U$, if:
$$\mathbf{x}(t) \in \mathcal{C} \Rightarrow \exists \mathbf{u} \in U \text{ such that } \phi(\mathbf{x}(t), \mathbf{u}(t)) \in \mathcal{C}, \text{ for all } t$$
- The set $\mathcal{C}_\infty \subseteq X$ is said to be the maximal control invariant set for the system $\mathbf{x}(t+1) = \phi(\mathbf{x}(t), \mathbf{u}(t))$ with constraints $\mathbf{x}(t) \in X$, $\mathbf{u}(t) \in U$, if it is control invariant and contains all control invariant sets contained in X
- These sets can be computed by using the MPT toolbox
<https://www.mpt3.org/>

Persistent feasibility lemma

- Define “truncated” feasibility set:

$$X_1 := \{\mathbf{x}_1 \in X \mid \exists (\mathbf{u}_1, \dots, \mathbf{u}_{N-1}) \text{ such that } \mathbf{x}_k \in X, \mathbf{u}_k \in U, k = 1, \dots, N-1, \\ \mathbf{x}_N \in X_f \text{ where } \mathbf{x}_{k+1} = A\mathbf{x}_k + B\mathbf{u}_k, k = 1, \dots, N-1\}$$

- Feasibility lemma: if set X_1 is a *control invariant set* for system:

$$\mathbf{x}(t+1) = A\mathbf{x}(t) + B\mathbf{u}(t), \quad \mathbf{x}(t) \in X, \quad \mathbf{u}(t) \in U, \quad t \geq 0$$

then the MPC law is persistently feasible

Persistent feasibility lemma

• Proof:

1. $\text{Pre}(X_1) = \{\mathbf{x} \in \mathbb{R}^n : \exists \mathbf{u} \in U \text{ such that } A\mathbf{x} + B\mathbf{u} \in X_1\}$

2. Since X_1 is control invariant

$$\forall \mathbf{x} \in X_1 \quad \exists \mathbf{u} \in U \text{ such that } A\mathbf{x} + B\mathbf{u} \in X_1$$

3. Thus $X_1 \subseteq \text{Pre}(X_1) \cap X$

4. One can write

$$X_0 = \{\mathbf{x}_0 \in X \mid \exists \mathbf{u}_0 \in U \text{ such that } A\mathbf{x}_0 + B\mathbf{u}_0 \in X_1\} = \text{Pre}(X_1) \cap X$$

5. Thus, $X_1 \subseteq X_0$

Persistent feasibility lemma

- Proof:

6. Pick some $\mathbf{x}_0 \in X_0$. Let U_0^* be the solution to the finite-time optimization problem, and \mathbf{u}_0^* be the first control. Let

$$\mathbf{x}_1 = A\mathbf{x}_0 + B\mathbf{u}_0^*$$

7. Since U_0^* is clearly feasible, one has $\mathbf{x}_1 \in X_1$. Since $X_1 \subseteq X_0$, one has

$$\mathbf{x}_1 \in X_0$$

hence the next optimization problem is feasible!

Practical significance

- For $N = 1$, we can set $X_f = X_1$. If we choose the terminal set to be control invariant, then MPC will be persistently feasible *independent* of chosen control objectives and parameters
- Designer can choose the parameters to affect performance (e.g., stability)
- How to extend this result to $N > 1$?

Persistent feasibility theorem

- Feasibility theorem: if set X_f is a *control invariant set* for system:

$$\mathbf{x}(t + 1) = A\mathbf{x}(t) + B\mathbf{u}(t), \quad \mathbf{x}(t) \in X, \quad \mathbf{u}(t) \in U, \quad t \geq 0$$

then the MPC law is persistently feasible

Persistent feasibility theorem

- Proof

1. Define “truncated” feasibility set at step $N - 1$:

$$X_{N-1} := \{\mathbf{x}_{N-1} \in X \mid \exists \mathbf{u}_{N-1} \text{ such that } \mathbf{x}_{N-1} \in X, \mathbf{u}_{N-1} \in U, \\ \mathbf{x}_N \in X_f \text{ where } \mathbf{x}_N = A\mathbf{x}_{N-1} + B\mathbf{u}_{N-1}\}$$

2. Due to the terminal constraint

$$A\mathbf{x}_{N-1} + B\mathbf{u}_{N-1} = \mathbf{x}_N \in X_f$$

Persistent feasibility theorem

- Proof

4. Since X_f is a control invariant set, there exists a $\mathbf{u} \in U$ such that

$$\mathbf{x}^+ = A\mathbf{x}_N + B\mathbf{u} \in X_f$$

5. The above is indeed the requirement to belong to set X_{N-1}

6. Thus, $A\mathbf{x}_{N-1} + B\mathbf{u}_{N-1} = \mathbf{x}_N \in X_{N-1}$

7. We have just proved that X_{N-1} is control invariant

8. Repeating this argument, one can recursively show that X_{N-2} , X_{N-3}, \dots, X_1 are control invariant, and the persistent feasibility lemma then applies

Practical aspects of persistent feasibility

- The terminal set X_f is introduced *artificially* for the sole purpose of leading to a *sufficient condition* for persistent feasibility
- We want it to be large so that it does not compromise closed-loop performance
- Though it is simplest to choose $X_f = 0$, this is generally undesirable
- We'll discuss better choices in the next lecture

Example

- System with $n = 3$ states, $m = 2$ inputs
- A, B chosen randomly
- quadratic stage cost: $c(\mathbf{x}, \mathbf{u}) = \mathbf{x}'\mathbf{x} + \mathbf{u}'\mathbf{u}$
- $X = \{\mathbf{x} \mid \|\mathbf{x}\|_\infty \leq 1\}$, $U = \{\mathbf{u} \mid \|\mathbf{u}\|_\infty \leq 0.5\}$
- initial point: $(0.9, -0.9, 0.9)$

Next time

- Stability of MPC, explicit MPC, and practical aspects